

AE 11: MLR Inference + conditions

Trail riders

Oct 26, 2022

! Important

The AE is due on GitHub by Saturday, October 29 at 11:59pm.

Packages

```
library(tidyverse)
library(tidymodels)
library(knitr)
```

Data

```
rail_trail <- read_csv("data/rail_trail.csv")
```

Exercise 1

Below is the model predicting volume from `hightemp` and `season`.

```
rt_mlr_main_fit <- linear_reg() %>%
  set_engine("lm") %>%
  fit(volume ~ hightemp + season, data = rail_trail)

tidy(rt_mlr_main_fit) |>
  kable(digits = 2)
```

| term | estimate | std.error | statistic | p.value |
|--------------|----------|-----------|-----------|---------|
| (Intercept) | -125.23 | 71.66 | -1.75 | 0.08 |
| hightemp | 7.54 | 1.17 | 6.43 | 0.00 |
| seasonSpring | 5.13 | 34.32 | 0.15 | 0.88 |
| seasonSummer | -76.84 | 47.71 | -1.61 | 0.11 |

Add an interaction effect between **hightemp** and **season** to the model. Do the data provide evidence of a significant interaction effect? Comment on the significance of the interaction terms.

```
interact <- linear_reg() %>%
  set_engine("lm") %>%
  fit(volume ~ hightemp + season + hightemp*season, data = rail_trail)

tidy(interact) %>%
  kable(digits = 3)
```

| term | estimate | std.error | statistic | p.value |
|-----------------------|----------|-----------|-----------|---------|
| (Intercept) | -10.530 | 166.796 | -0.063 | 0.950 |
| hightemp | 5.482 | 2.950 | 1.858 | 0.067 |
| seasonSpring | -293.949 | 190.331 | -1.544 | 0.126 |
| seasonSummer | 354.179 | 255.078 | 1.389 | 0.169 |
| hightemp:seasonSpring | 4.884 | 3.262 | 1.497 | 0.138 |
| hightemp:seasonSummer | -4.537 | 3.754 | -1.209 | 0.230 |

These data, surprisingly, do not provide evidence of a significant interaction effect. The p-value of **hightemp * Spring** and **Summer** are .138 and .230, respectively.

```
glance(rt_mlr_main_fit) %>%
  select(adj.r.squared, AIC, BIC)
```

```
# A tibble: 1 x 3
  adj.r.squared  AIC    BIC
    <dbl> <dbl> <dbl>
1      0.365 1093. 1106.
```

```
glance(interact) %>%
  select(adj.r.squared, AIC, BIC)
```

```
# A tibble: 1 x 3
  adj.r.squared  AIC    BIC
      <dbl> <dbl> <dbl>
1      0.435 1084. 1102.
```

Exercise 2

Below is the model predicting volume from all available predictors.

```
rt_full_fit <- linear_reg() %>%
  set_engine("lm") %>%
  fit(volume ~ ., data = rail_train)

tidy(rt_full_fit) |>
kable(digits = 2)
```

| term | estimate | std.error | statistic | p.value |
|-----------------|----------|-----------|-----------|---------|
| (Intercept) | 17.62 | 76.58 | 0.23 | 0.82 |
| hightemp | 7.07 | 2.42 | 2.92 | 0.00 |
| avgtemp | -2.04 | 3.14 | -0.65 | 0.52 |
| seasonSpring | 35.91 | 32.99 | 1.09 | 0.28 |
| seasonSummer | 24.15 | 52.81 | 0.46 | 0.65 |
| cloudcover | -7.25 | 3.84 | -1.89 | 0.06 |
| precip | -95.70 | 42.57 | -2.25 | 0.03 |
| day_typeWeekend | 35.90 | 22.43 | 1.60 | 0.11 |

Fill in the code to plot the histogram of residuals with an overlay of the normal distribution based on the results of the model.


i Note

Update to `eval: true` once the code is updated.

```
rt_full_aug <- augment(rt_full_fit$fit)

ggplot(rt_full_aug, aes(.resid)) +
  geom_histogram(aes(y = after_stat(density)), binwidth = 50) +
  stat_function(
    fun = dnorm,
    args = list(mean = mean(rt_full_aug$.resid), sd = sd(rt_full_aug$.resid)),
```

```
lwd = 2,  
color = "red"  
)
```



Submission

! Important

To submit the AE:

- Render the document to produce the PDF with all of your work from today's class.
- Push all your work to your `ae-11-` repo on GitHub. (You do not submit AEs on Gradescope).